

Methods

In this study, we conducted two simulations to evaluate the performance of LASSO models in identifying important features in high-dimensional data. The first simulation assumes sparsity of the features in the pixel space, while the second simulation assumes sparsity in the frequency space.

The pixel space refers to the original high-dimensional space where each dimension represents a pixel in an image. In this space, features are directly observed and may have inherent correlations. Conversely, the frequency space is a transformed version of the pixel space, obtained through techniques like eigen decomposition.

Suppose X is a column vector representing 256 pixels. Its covariance matrix, Σ , is defined to have an exponential correlation structure, where $\Sigma_{ij} = -\exp(\text{dist}(i, j))$. Here, $\text{dist}(i, j)$ is the distance between the pixels i and j in a 16×16 matrix.

Let V be the matrix of eigenvectors of Σ , with each column representing an eigenvector. We can transform the random vector X into the frequency space by $X_{\text{freq}} = V^T X$. The covariance matrix of X_{freq} is given by $\text{cov}(X_{\text{freq}}) = V^T \Sigma V$, which is a diagonal matrix.

For the simulations, in each iteration, we randomly generate X_{freq} from a multivariate normal distribution with the covariance matrix $\text{cov}(X_{\text{freq}})$. We repeat this process 1000 times. Then, we calculate X as $X = V X_{\text{freq}}$.

In the first simulation, we assume sparsity in the coefficient vectors in the pixel space. The coefficient vector β was specified to have non-zero values exclusively within a central 8×8 region. The response variable y was drawn from a binomial distribution with success probabilities determined by $\eta = X\beta$. The non-zero coefficients in β were chosen such that the probability $p = \frac{1}{1+\exp(-\eta)}$ was uniformly distributed across interval $[0, 1]$.

In the second simulation, we assume sparsity in the coefficient vectors in the frequency space. We defined a sparse coefficient vector b in the frequency space, where most of the 256 entries were zero and a randomly 10% were non-zero. The response variable y was generated similarly to the first simulation, ensuring $p = \frac{1}{1+\exp(-\eta)}$ was evenly distributed.

For both simulations, we fit two models: one using the covariates in the pixel space and another using the covariates in the frequency space. Each dataset, generated in size 1000×256 and representing images of 16×16 pixels, was split into training (80%) and test (20%) sets. The regularization parameter λ was tuned using cross-validation with the default binomial deviance metric. The dataset was divided into 10 folds, with the model trained and validated iteratively across these folds, varying λ . Here we use two λ values, ‘lambda.min’ is the regularization value that minimizes cross-validated error, and ‘lambda.1se’ is the largest lambda within one standard error of the minimum.

After selecting the optimal λ , model performance was evaluated using accuracy and AUC metrics. Additionally, a permutation test was conducted 100 times to calculate p-values for each covariate. Across all iterations, we calculated the mean and standard deviation of the metrics, as well as the percentage of significant p-values for each covariate.

Results

In Simulation 1, the distribution of the success probability p was evaluated at various β values: 0.01, 0.05, 0.1, 0.2, and 1. Figure 1 shows that $\beta = 0.1$ resulted in the most even spread of p , making it the optimal choice for model fitting.

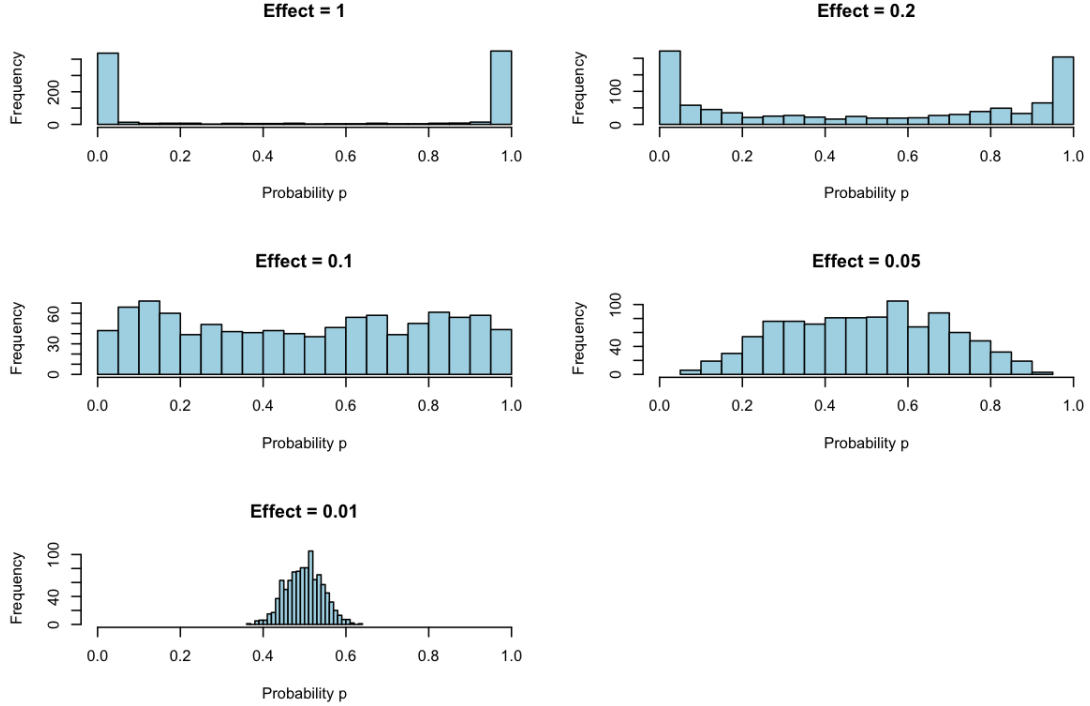


Figure 1: The distribution of p at different β values.

In Simulation 2, similarly, the distribution of p was evaluated at various b values: 0.1, 0.2, 0.4, 0.6, 0.8, and 1. Figure 2 shows that $b = 0.4$ yielded the most evenly distributed p .

Using the chosen effects, we drew the response variable y from a binomial distribution with success probabilities determined by $\eta = X\beta$ in both simulations. We then plotted the group mean difference in the covariate values between instances where $y = 1$ and $y = 0$ in both the pixel space and the frequency space. Figure 3

Figure 3 shows the group mean difference in the covariate values between instances and the frequency space. The heatmaps illustrate that in the pixel space, the central region with non-zero coefficients in β corresponds to higher mean values of the covariates. Indicating its strong association with the positive response. In the frequency space, the pattern is more dispersed but still reflects the influence of the non-zero coefficients.

Figure 4 presents the coefficients for Simulation 1, where sparsity was assumed in the pixel space. The coefficients exhibit a concentrated non-zero region in the center of the image, reflecting the assumption that only a central subset of the covariates significantly influences the response variable. This central region aligns with the higher mean values observed in the pixel space's covariates, as shown in Figure 3.

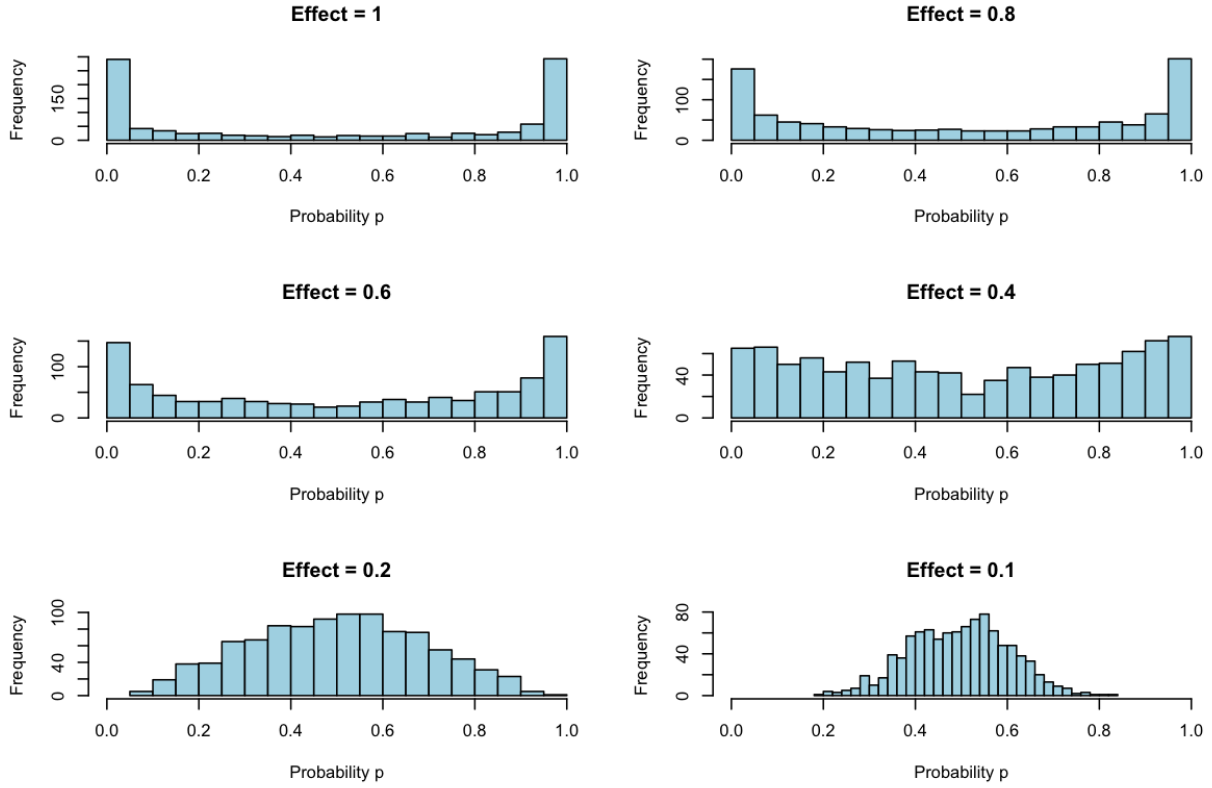


Figure 2: The distribution of p at different b values. $b = 0.4$ was chosen for model fitting as it gives the most evenly distributed values.




Figure 3: The group mean difference between $y = 1$ and $y = 0$ in the pixel space and frequency space.

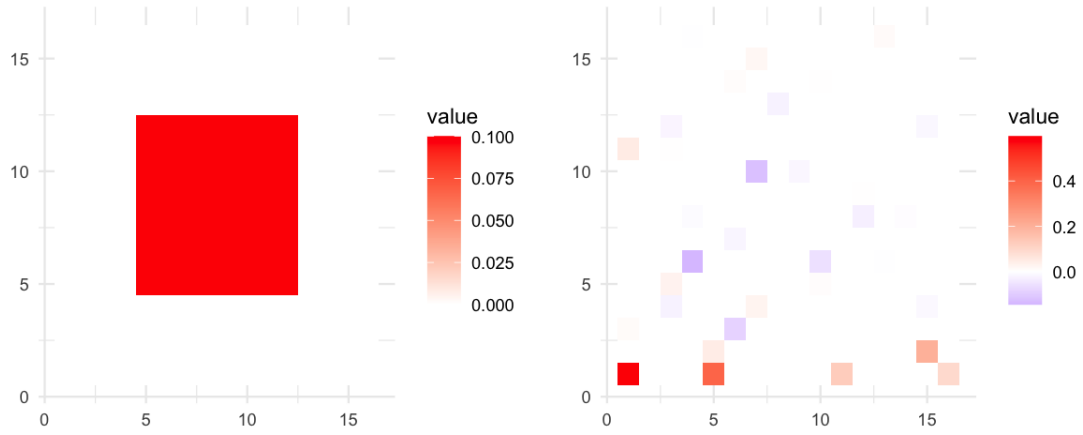


Figure 4: Coefficients for Simulation 1 in the pixel space (left) and frequency space (right).




Figure 5: Coefficients for Simulation 2.




Figure 6: The group mean difference between $y = 1$ and $y = 0$.

Table 1 shows the average AUCs and accuracys calculated by fitting models on the 1000 simulated dataset, on both the frequency sapce and pixel space. We can see that no matter we assume sparsity in coefficient vectors in the pixel space of the frequency space, fitting models on the frequency space provides better performance in AUCs and accuracys. When assuming sparsity in coefficient vectors in the pixel space and using ‘lambda.min’ as the regularization value, models fitting with covariates in the pixel space (X) achieves an AUC of 0.803 (se = 0.031), and an accuracy of 72.5% (se = 0.032); while models fitting with covariates in the corresponding frequency space (X_{freq}) achieves a slightly better AUC of 0.827 (se = 0.029) and better accuracy of 74.6% (se = 0.031). We see a similar advantage of fitting models on the frequency space when assuming sparsity in coefficient vectors in the frequency space, and using either ‘lambda.min’ or ‘lambda.1se’ does not affect this advantage.

Table 1:

Lambda	Model on the pixel space		Model on the freq space	
	AUC (SE)	Accuracy (SE)	AUC (SE)	Accuracy (SE)
lambda.min	0.803 (0.031)	0.725 (0.032)	0.827 (0.029)	0.746 (0.031)
lambda.1se	0.800 (0.031)	0.723 (0.032)	0.826 (0.029)	0.745 (0.031)
lambda.min	0.756 (0.036)	0.687 (0.035)	0.816 (0.032)	0.735 (0.033)
lambda.1se	0.734 (0.040)	0.669 (0.036)	0.816 (0.032)	0.734 (0.034)

Then, for both Simulation 1 and Simulation 2, we fitted two models, one using covariates on the frequency space, another using covariates on the pixel space. In

Figure 7 shows the frequencies tied to the top three eigenvalues representing the main patterns in the pixel space. The frequency associated with the smallest eigenvalue highlights the least significant variance.




Figure 7: Frequencies associated with the top three eigenvalues (top row and bottom left) and the frequency associated with the smallest eigenvalue (bottom right), highlighting the primary and least significant patterns in the pixel space.